

🎯 Reinforcement Learning (RL) 🎯

<https://drive.google.com/file/d/1gWmitl1j7mam4ZmdYC5csS5JW7voF33-/view?pli=1>



What is Reinforcement Learning?

Reinforcement Learning (RL) is a type of Machine Learning where an **agent learns by interacting with an environment**.

The agent:

- Performs actions,
- receives rewards or penalties,
- learns from experience,
- and improves its decisions over time.

The main idea is: **“Learn which actions give the highest rewards.”**

RL solves **sequential decision-making problems**.

This means:

- actions are taken **one after another**
- each **action** affects the **future state**
- the **agent must decide the best sequence of actions to reach a goal**

◆ Goal of RL

The objective is: **To learn to act optimally in the environment**

The agent should learn:

- what action to take,

- in which state,
- to get maximum reward.

◆ Rewards and Feedback

After every action, the **agent receives feedback** called a **reward**.

- *Positive reward → good action*
- *Negative reward (penalty) → bad action*

The agent **learns from these rewards**

◆ Utility of Agent

- **utility (performance)** of an RL agent is **defined by the rewards it collects**.
- **More rewards = better performance.**
- So the agent must: **maximize expected rewards**

◆ Learning from Samples

- RL learns through **trial and error**.
- The agent:
 - observes outcomes,
 - stores experiences,
 - improves future decisions.

It **does NOT need predefined correct answers**.

Components of Reinforcement Learning

Component	Meaning
-----------	---------

Agent	Learner/decision maker
Environment	World where agent operates
State (S)	Current situation
Action (A)	Move taken by agent
Reward (R)	Feedback received
Policy (π)	Strategy of choosing actions

Applications of Reinforcement Learning

 **Game Playing:** Agents learn winning strategies by playing repeatedly.

 **Self-Driving Cars**

RL helps cars: make driving decisions, avoid collisions, choose safe routes.

 **Robotics**

Robots learn: Walking, object picking, Balancing, navigation.

 **Stock Market & Trading**

RL agents learn when to buy or sell, how to maximize profit.

 **Healthcare**

- treatment planning, medicine dosage optimization, patient monitoring.

 **Recommendation Systems**

- Netflix etc use RL to improve recommendations.
-



Types of Reinforcement Learning

1 Positive Reinforcement

- Good behavior is rewarded.
- **Example**
 - Robot reaches destination: Reward = +10
 - This increases chances of repeating the action.

2 Negative Reinforcement

- Undesirable condition is removed after correct action.
- **Example**
 - Alarm stops when correct action is taken.
 - Encourages good behavior.

3 Model-Based RL

- Agent has a model of environment.
- It can:
 - *predict outcomes,*
 - *plan future actions.*

4 Model-Free RL

- Agent learns only from experience.
 - No environment model is available.
 - Example: **Q-Learning**
-




Markov Decision Process (MDP)

- RL problems are usually modeled using MDP
- MDP provides a **mathematical framework** for RL.

An MDP is defined by:

Component	Meaning
States (S)	Possible situations
Actions (A)	Possible actions
Transition Function $T(s,a,s')$	Probability of moving to next state
Reward Function $R(s,a,s')$	Reward received after transition

Sometimes **Start state (s_0)** & **End/terminal state** are also included.



What Does “Markov” Mean?

- “Markov” means: **future depends only on the present state.**
- The **past does NOT** matter if current state is known.

Markov Property

Instead of depending on all previous states: $P(s_{t+1} \mid s_t, a_t, s_{t-1}, a_{t-1}, \dots)$

the next state depends only on:

$$P(s_{t+1} \mid s_t, a_t)$$

Meaning: **Current state + current action** completely determine future transition.

Policy in RL $\pi(s)$

- **Policy** is a **recommendation function**
 - tells the agent **what action to take in a given state.**
 - It is represented as: $\pi(s)$
 - **Input** → current state
 - **Output** → action to take
 - The **best policy** is called: **optimal policy** π^*
 - It is the policy that gives: maximum rewards, maximum utility, best long-term outcome.
 - For MDPs, we want to achieve the optimal policy
-

Learning Optimality

- The goal of RL is: **Collect maximum rewards**
- The agent **tries to maximize total reward over time.**
- 'The optimality criteria is to maximize the rewards'

Total Reward (Undiscounted Utility)

If:

r_t = reward at time t

then total reward is:

$$r_0 + r_1 + r_2 \dots = \sum_{t=0}^{\infty} r_t$$

(Undiscounted utility)

> All rewards are **treated equally.**

Discounted Utility

- In RL, **future rewards** are usually **considered less important.**
- So we multiply future rewards by a **discount factor** [**gamma**]:

$$0 \leq \gamma \leq 1$$

$$r_0 + \gamma r_1 + \gamma^2 r_2 \dots = \sum_{t=0}^{\infty} \gamma^t r_t$$

Rewards farther in future become **highly discounted.**

γ Value	Meaning
$\gamma = 0$	Only immediate reward matters
γ close to 1	Future rewards also important

★ Value Functions

- **Value Function** tells us: “How good or useful something is.”
- It measures the **expected reward** an agent can receive in the future.
- **utility(usefulness) of a sequence of rewards** is defined as a value function

Value functions help the agent:


- evaluate states,
- evaluate actions,
- choose better decisions,
- maximize rewards.

● Types of Value Functions

Type	Meaning
State Value Function $V(s)$	Goodness of being in a state
Action Value Function $Q(s,a) / V(s,a)$	Goodness of taking action (a) in state (s)

📌 State Value Function — $V(s)$

- A state value tells: “How good is it to be in this state?”
- It estimates the **future rewards expected from that state.**




 **Interpretation:** Higher value = better chance of receiving rewards.

State	Value	Meaning
(s_0)	0.93	Good state
(s_1)	0.94	Slightly better
(s_2)	0.95	Even better

Color	Meaning
Green	Positive/normal states
Red	Negative reward state
Gray	Obstacle/block
Bright Green	Goal state (+1 reward)

🔥 Q-Values (Action Value Function)

- A Q-value tells: **“How good is it to take action (a) from state (s)?”**
- Q-values **evaluate** specific actions.



Action	Q-Value	Meaning
Up	0.59	Good action
Right	0.59	Good action
Down	0.53	Less useful
Left	0.53	Less useful

The agent always prefers: **actions with higher Q-values.**

Because: **they lead to better long-term rewards.**

🤖 Reinforcement Learning Methods

◆ Unknown Model in RL

- In Reinforcement Learning: **The environment model is usually unknown.**
- Specifically:
 - Transition function (T)
 - Reward function (R)

are **NOT** given.

- The agent **must** learn through interaction.
-

RL Methods Categories

RL methods mainly fall into:

Type	Approach
Model-Based RL	Learn model first
Model-Free RL	Learn directly from experience

Model-Based Reinforcement Learning

First learn:

- transition function (T)
- reward function (R)

by interacting with environment.

Then use them to: **compute optimal policy.** [thru Dynamic programming]

2 Model-Free Reinforcement Learning

The agent directly learns:

- value functions,
- policies,

WITHOUT learning the environment model.

› **Example:** Q - learning

 **Learns From**

- trial and error,
 - rewards,
 - experience.
-

Q-Learning Algorithm

Q-Learning is a:

- **Model-free** Reinforcement Learning algorithm
- **Online learning** algorithm
- based on **trial and error**

It learns: “Which action is best in a particular state.”

› Q-Learning **improves the estimates of Q-values** using:

1. *Immediate reward*
2. *Future discounted rewards*

The agent **continuously updates its knowledge after every action.** [Temporal difference learning]

› Q-Learning uses a concept called: **Temporal Difference Learning**

This means:

- learning happens **immediately after experience**,
- estimates are updated using **current reward + future estimate**.

Instead of waiting until the end: the agent learns step-by-step.

Agent:

1. observes current state
2. chooses action
3. performs action
4. receives reward
5. moves to next state
6. updates Q-value

This repeats continuously.

Action Selection Using Policy

The agent chooses action using:

$$\pi(s) = \arg \max Q(s, a)$$

- The **policy selects** the action with the **highest Q-value**.
- agent chooses: **the action expected to give maximum reward**.

Transition in Q-Learning

Suppose the agent:

- **starts in state (s)**
- **takes action (a)**
- **reaches next state (s')**

During this transition: it **receives reward (r)**

From this experience, the agent discovers:

- ☒ **Immediate reward** r
- ☒ **Next state's Q-values** $Q(s', a)$

This information is used to improve learning.

Sample Estimate

The sampled experience is:

$$\text{Sample} = r + \gamma \max_a Q(s', a)$$

Symbol	Meaning
r	Immediate reward
gamma	Discount factor
$\max Q(s', a)$	Best future reward

new estimate should depend on: **current reward** PLUS **best possible future reward**.

🔥 Q-Learning Update Rule

$$Q(s, a) = Q(s, a) + \alpha[r + \gamma \max_a Q(s', a) - Q(s, a)]$$

› adds the **sample to** the **old estimate** of Q-value using a **learning rate α** as running average

› It updates old knowledge using:

- current experience,
- future expectation

$$0 \leq \alpha \leq 1$$

High α **Learn quickly**

Low α **Learn slowly**

↺ Running Average Interpretation

$$Q(s, a) = \alpha(\text{sample}) + (1 - \alpha)Q(s, a)$$

- part old knowledge,
- part new experience.

🧩 Q-Learning Algorithm Steps

Algorithm 1 Q-learning as proposed by Watkins and Dayan (1992)

Require: discount factor γ , learning rate α

```
1: Initialize  $Q(s,a)$  arbitrarily
2: for all episodes do
3:   Initialize starting state  $s$ 
4:   repeat
5:     Choose action  $a$  using  $Q$  and an exploration strategy
6:     Execute  $a$ 
7:     Observe next state  $s'$  and receive immediate reward  $r$ 
8:     Update,  $Q(s,a) \leftarrow Q(s,a) + \alpha \left( r + \gamma \max_{a'} Q(s',a') - Q(s,a) \right)$ 
9:      $s := s'$ 
10:  until  $s'$  is not the terminal state
11: end for
```

Step 1 Initialize Q-values:

- Initialize **$Q(s,a)$** randomly or with zeros.

Step 2 Start Episode

- Choose **starting state**: s

Step 3 Choose Action

Select action:

- **using current Q-values,**
- **plus exploration strategy.**

Step 4 Execute Action

- **Take action a**

Step 5 Observe Environment

Observe:

- **next state (s')**
- **reward (r)**

Step 6 Update Q-value

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_a Q(s', a) - Q(s, a)]$$

Step 7 Move to Next State

- Update: $s = s'$

Step 8 Repeat

Continue until:

- **terminal/goal state reached.**
-



Why Q-Learning Works

Over time:

- **good actions receive high Q-values,**
- **bad** actions receive **low** Q-values.

Eventually: **optimal path emerges.**

- Q-Learning is: **Off-Policy RL**
 - it **learns the optimal policy, regardless of the actions currently taken.**
- Even random exploration can still lead to optimal learning.
- The algorithm tries to learn: $Q^*(s, a)$ which is the **optimal Q-value.**
- The optimal policy is: $\pi^*(s) = \arg \max_a Q^*(s, a)$
- The best action is: **the one with highest optimal Q-value.**

Q-Learning **converges to optimal policy** if:

1 Infinite Visits

- **Every state-action pair** is **visited many times**
- **ensures proper learning.**

2 Learning Rate Decays

- **Learning rate** should **gradually decrease over time.**

Conditions:

$$\sum_{t=0}^{\infty} \alpha_t = \infty \quad \& \quad \sum_{t=0}^{\infty} \alpha_t^2 < \infty$$

These mathematical conditions **guarantee convergence.**

In practice:

- **policies often converge much earlier**, before exact Q-values fully converge.

This makes Q-learning:

- practical,
- efficient,
- widely usable.



Exploration Strategy

- If the **agent always chooses: current best action**, it **may miss: better unseen paths**. So exploration is necessary.

Exploration	Try new actions
Exploitation	Use best known action

RL balances both

★ Epsilon-Greedy Strategy

Most common exploration method: **ϵ -Greedy Policy**

Probability	Action
ϵ	Random action
$(1 - \epsilon)$	Best Q-value action

If: epsilon = 0.1

Then:

- **10%** → *random exploration*
- **90%** → *best known action*

🧠 Why ϵ -Greedy Works

It allows:

- discovering better states,
- avoiding local optimum,
- visiting more state-action pairs.

📖 Complete Intuition of Q-Learning

The agent repeatedly asks:

- "Was this action good?"
- "Could future rewards make it even better?"

Then it updates its knowledge accordingly.


Over time:

- optimal actions dominate,
- poor actions disappear.



Q-Learning Example — Room Navigation Problem

<https://www.youtube.com/watch?v=J3qX50yyiU0>

- Suppose we have 5 rooms in a building connected by doors as shown in the figure below. We'll number each room 0 through 4. The outside of the building can be thought of as one big room (5). Notice that doors 1 and 4 lead into the building from room 5 (outside).



- We can represent the rooms on a graph, each room as a node, and each door as a link.



$$R = \begin{matrix} & \text{Action} \\ \text{State} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} -1 & -1 & -1 & -1 & 0 & -1 \\ -1 & -1 & -1 & 0 & -1 & 100 \\ -1 & -1 & -1 & 0 & -1 & -1 \\ -1 & 0 & 0 & -1 & 0 & -1 \\ 0 & -1 & -1 & 0 & -1 & 100 \\ -1 & 0 & -1 & -1 & 0 & 100 \end{bmatrix} \end{matrix}$$

> Reward Matrix R:

This matrix tells us: “What reward do we get for moving from one room to another?”

> basically a map of: allowed moves, blocked moves, goal rewards.

- $R[1,5]=100 \Rightarrow$ Going from room 1 $\rightarrow$ room 5 gives reward 100.

-1	Invalid move
0	Valid move but no reward
100	Goal reached

Q learning algorithm

For each s, a initialize the table entry $\hat{Q}(s, a)$ to zero.

Observe the current state s

Do forever:

- Select an action a and execute it
- Receive immediate reward r
- Observe the new state s'
- Update the table entry for $\hat{Q}(s, a)$ as follows:

$$\hat{Q}(s, a) \leftarrow r + \gamma \max_{a'} \hat{Q}(s', a')$$

- $s \leftarrow s'$

- Learning rate = 0.8 and the initial state as Room 1.
- Initialize matrix Q as a zero matrix:

> Q Matrix:

- "How useful is moving from one room to another?"
- Initially 0 $\Rightarrow$ because the agent knows nothing.

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$

$$Q(s, a) = R(s, a) + \gamma \max_{a'} Q(s', a')$$

- Look at the second row (state 1) of matrix R.
- There are two possible actions for the current state 1: go to state 3, or go to state 5.
- By random selection, we select to go to 5 as our action.

> if we select 5 $\Rightarrow$ we get **immediate reward of 100**

- Now let's imagine what would happen if our agent were in state 5 (next state).
- Look at the sixth row of the reward matrix R (i.e. state 5).
- It has 3 possible actions: go to state 1, 4 or 5.

- $Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \text{Max}[Q(\text{next state}, \text{all actions})]$
- $Q(1, 5) = R(1, 5) + 0.8 * \text{Max}[Q(5, 1), Q(5, 4), Q(5, 5)] = 100 + 0.8 * 0 = 100$

		Action												
State		0	1	2	3	4	5							
$R =$	0	-1	-1	-1	-1	0	-1	$Q =$	0	0	0	0	0	0
	1	-1	-1	-1	0	-1	100		1	0	0	0	0	0
	2	-1	-1	-1	0	-1	-1		2	0	0	0	0	0
	3	-1	0	0	-1	0	-1		3	0	0	0	0	0
	4	0	-1	-1	0	-1	100		4	0	0	0	0	0
	5	-1	0	-1	-1	0	100		5	0	0	0	0	0

- The next state, 5, now becomes the current state.
- Because 5 is the goal state, we've finished one episode.
- Our agent's brain now contains an updated matrix Q as:

		0	1	2	3	4	5
$Q =$	0	0	0	0	0	0	0
	1	0	0	0	0	0	100
	2	0	0	0	0	0	0
	3	0	0	0	0	0	0
	4	0	0	0	0	0	0
	5	0	0	0	0	0	0

- For the next episode, we randomly choose the initial state – say 3 (can go to 1, 2 & 4)

		Action					
State		0	1	2	3	4	5
$R =$	0	-1	-1	-1	-1	0	-1
	1	-1	-1	-1	0	-1	100
	2	-1	-1	-1	0	-1	-1
	3	-1	0	0	-1	0	-1
	4	0	-1	-1	0	-1	100
	5	-1	0	-1	-1	0	100

- Now we imagine that we are in state 1 (next state).

- Now we imagine that we are in state 1 (next state).
- Look at the second row of reward matrix R (i.e. state 1).
- It has 2 possible actions: go to state 3 or state 5.
- Then, we compute the Q value:
- $Q(\text{state}, \text{action}) = R(\text{state}, \text{action}) + \text{Gamma} * \text{Max}[Q(\text{next state}, \text{all actions})]$
- $Q(3, 1) = R(3, 1) + 0.8 * \text{Max}[Q(1, 3), Q(1, 5)] = 0 + 0.8 * \text{Max}(0, 100) = 80$

State	Action					
	0	1	2	3	4	5
0	-1	-1	-1	-1	0	-1
1	-1	-1	-1	0	-1	100
2	-1	-1	-1	0	-1	-1
3	-1	0	0	-1	0	-1
4	0	-1	-1	0	-1	100
5	-1	0	-1	-1	0	100

State	Action					
	0	1	2	3	4	5
0	0	0	0	0	0	0
1	0	0	0	0	0	0
2	0	0	0	0	0	100
3	0	80	0	0	0	0
4	0	0	0	0	0	0
5	0	0	0	0	0	0

> repeat this baar baar, n we get our Q matrix

- If our agent learns more through further episodes, it will finally reach convergence values in matrix Q like:

State	Action					
	0	1	2	3	4	5
0	0	0	0	0	80	0
1	0	0	0	64	0	100
2	0	0	0	64	0	0
3	0	80	51	0	80	0
4	64	0	0	64	0	100
5	0	80	0	0	80	100

- Tracing the best sequences of states is as simple as following the links with the highest values at each state.

