

Chapter 14: Vector Space Classification

<https://drive.google.com/file/d/1sWvloE15gwzy-4MfuBntBKjSYAKbr2Fv/view>

https://drive.google.com/file/d/1uB9leiG-HEAtEX9IgSXdg_DFDWvUVAT4/view

1. Vector Space Model for Classification

- **each document** is represented as a **feature vector** in a high-dimensional space (**one dimension per vocabulary term**).



Contiguity Hypothesis:

- > **Documents in the same class form a contiguous region in vector space**, and **regions of different classes do not overlap**.
- > This **allows a classifier to draw decision boundaries**.

> **Classification Pipeline**

- **Corpus** – the classification dataset
- **Pre-processing** – tokenisation, stemming, stopword removal
- **Feature selection & weighting** – TF-IDF, MI, χ^2
- **Similarity/distance function**
- **Classification algorithm**
- **Performance evaluation**

2. Rocchio's Algorithm

- Rocchio's algorithm adapts the relevance feedback technique for classification.
- It **divides** the **vector space into regions centred on class centroids** (prototypes)
- simple and efficient
- **inaccurate** if **classes** are **not approximately spheres** with similar radii.

› Algorithm:

```
TRAINROCCHIO( $\mathcal{C}, \mathbb{D}$ )  
1  for each  $c_j \in \mathcal{C}$   
2  do  $D_j \leftarrow \{d : \langle d, c_j \rangle \in \mathbb{D}\}$   
3      $\vec{\mu}_j \leftarrow \frac{1}{|D_j|} \sum_{d \in D_j} \vec{v}(d)$   
4  return  $\{\vec{\mu}_1, \dots, \vec{\mu}_J\}$ 
```

```
APPLYROCCHIO( $\{\vec{\mu}_1, \dots, \vec{\mu}_J\}, d$ )  
1  return  $\arg \min_j |\vec{\mu}_j - \vec{v}(d)|$ 
```

For each class c_j , compute the **centroid** (centre of mass) of all training documents in that class:

$$\vec{\mu}_j \leftarrow \frac{1}{|D_j|} \sum_{d \in D_j} \vec{v}(d)$$

Assign a new document d to the class whose **centroid is closest**:

$$\arg \min_j |\vec{\mu}_j - \vec{v}(d)|$$

› Worked Example – Rocchio

Doc	Words	Class	Vector (beijing, chinese, japan, macao, shanghai, tokyo)
1	Chinese Beijing Chinese	c	$\langle 1, 2, 0, 0, 0, 0 \rangle$
2	Chinese Chinese Shanghai	c	$\langle 0, 2, 0, 0, 1, 0 \rangle$
3	Chinese Macao	c	$\langle 0, 1, 0, 1, 0, 0 \rangle$
4	Tokyo Japan Chinese	j	$\langle 0, 1, 1, 0, 0, 1 \rangle$
5 (test)	Chinese Chinese Chinese Tokyo Japan	?	$\langle 0, 3, 0, 0, 0, 1 \rangle$

- $\mu^c = (1/3)(d1 + d2 + d3) = \langle 1/3, 5/3, 0, 1/3, 1/3, 0 \rangle$
- $\mu^j = d4 = \langle 0, 1, 1, 0, 0, 1 \rangle$
- $\text{Distance}(d5, \mu^c) = 5$

✓ Distance to class c

$$d(doc5, \mu_c)$$

Compute difference:

$$\begin{aligned} & (0 - 1/3, 3 - 5/3, 0 - 0, 0 - 1/3, 0 - 1/3, 1 - 0) \\ & = (-1/3, 4/3, 0, -1/3, -1/3, 1) \end{aligned}$$

Square & sum:

$$= (1/9 + 16/9 + 0 + 1/9 + 1/9 + 1) \approx 5$$

- $\text{Distance}(d5, \mu^j) = 4$

Result: d5 is classified as class j (closer to j centroid)

› Weaknesses of Rocchio Classification

- **Inaccurate** when classes are not **approximately spherical** in shape
- Fails when classes have **similar radii but are non-convex or multimodal**
- A document **can only be assigned to one class** (single prototype per class)

3. k-Nearest Neighbour (kNN) Classification

- kNN is an **instance-based (lazy) learning method**.
- It **does not build an explicit model during training** – instead it **memorises** the **entire training set** and **classifies at query time**.
- **Memory based method**

› Algorithm

Training: Simply **store all training examples and their feature vectors**. No model is built.

Testing for document d:

- Compute **cosine similarity between d** and **all training documents**
- **Find** the **k most similar** (nearest) **neighbours**
- **Return** the **majority class** among those k neighbours

› kNN Example

	Doc	Words	Class
Training	1	Chinese Beijing Chinese	c
	2	Chinese Chinese Shanghai	c
	3	Chinese Macao	c
	4	Tokyo Japan Chinese	j
Test	5	Chinese Chinese Chinese Tokyo Japan	?

Dictionary:

⟨beijing, chinese, japan, macao, shanghai, tokyo⟩

Vectors:

Doc	Vector	Class
d1	⟨1,2,0,0,0,0⟩	c
d2	⟨0,2,0,0,1,0⟩	c
d3	⟨0,1,0,1,0,0⟩	c
d4	⟨0,1,1,0,0,1⟩	j
d5	⟨0,3,0,0,0,1⟩	? (test)

Using the same Chinese/Japan dataset:

$$\cos(\theta) = \frac{d_i \cdot d_j}{\|d_i\| \|d_j\|}$$

- $\text{Cos}(d1, d5) = 0.848 \rightarrow \text{class c}$
- $\text{Cos}(d2, d5) = 0.848 \rightarrow \text{class c}$
- $\text{Cos}(d3, d5) = 0.424 \rightarrow \text{class c}$
- $\text{Cos}(d4, d5) = 0.953 \rightarrow \text{class j}$
- **1-NN** (k=1): d5 → **class j** (most similar single neighbour is d4)
- **3-NN** (k=3): d5 → **class c** (majority of top 3 are class c: d1, d2, d4 → 2c vs 1j, but d4 is the 1st, d1 & d2 also top 3)

Advantages of kNN

- Simple, intuitive, and easy to implement
- **Only one hyperparameter (k) to tune**
- No assumptions about data distribution
- Naturally handles multiclass problems

Disadvantages of kNN

- **Sensitive to the choice of k** , distance function, and noisy data
- **Lazy learner** – no precomputed model (all computation at test time)
- **No training cost**
- **very large computation and storage cost** for large datasets
- All **features treated equally** (noisy features degrade performance)

◀ Food for Thoughts ▶

Q4. Rocchio's Algorithm for Text Classification

Rocchio's algorithm computes a prototype (centroid) for each class from the training documents and classifies new documents by finding the closest centroid.

Modification: The algorithm can be adapted by weighting training documents (e.g. by TF-IDF) before computing centroids, which gives more important terms more influence in the centroid computation.

Q5. Weaknesses of Rocchio Classification

- Assumes classes are approximately spherical – fails for non-convex or irregular shapes
- One centroid per class cannot capture multimodal distributions
- Sensitive to outliers which distort the centroid
- Cannot handle overlapping classes well

Q6. Benefits of kNN for Text Classification

- Simple and easy to implement – no training phase required
- No assumption about class shape or distribution
- Naturally handles multiclass problems
- Can be updated by simply adding new training examples
- Works well with sufficient training data

Q7. Multinomial NB – Classification Example

Split	Doc ID	Words in Document	In c = China?
Training	1	Taipei Taiwan	yes
Training	2	Macao Taiwan Shanghai	yes
Training	3	Japan Sapporo	no
Training	4	Sapporo Osaka Taiwan	no
Test	5	Taiwan Taiwan Sapporo	?

Part (a): Class Prior Probabilities

- $P(\text{China} = \text{yes}) = 2/4 = 0.5$
- $P(\text{China} = \text{no}) = 2/4 = 0.5$

Part (b): Conditional Probabilities (Multinomial NB with Laplace smoothing)

Vocabulary $V = \{\text{Taipei, Taiwan, Macao, Shanghai, Japan, Sapporo, Osaka}\} \rightarrow |V| = 7$

Class YES (China): Words are Taipei, Taiwan, Macao, Taiwan, Shanghai → total tokens = 5

- $P(\text{Taiwan} | \text{yes}) = (2+1)/(5+7) = 3/12 = 1/4$
- $P(\text{Sapporo} | \text{yes}) = (0+1)/(5+7) = 1/12$
- $P(\text{Taipei} | \text{yes}) = (1+1)/(5+7) = 2/12 = 1/6$
- $P(\text{Macao} | \text{yes}) = (1+1)/(5+7) = 2/12 = 1/6$
- $P(\text{Shanghai} | \text{yes}) = (1+1)/(5+7) = 2/12 = 1/6$
- $P(\text{Japan} | \text{yes}) = (0+1)/(5+7) = 1/12$
- $P(\text{Osaka} | \text{yes}) = (0+1)/(5+7) = 1/12$

Class NO: Words are Japan, Sapporo, Sapporo, Osaka, Taiwan → total tokens = 5

- $P(\text{Taiwan} | \text{no}) = (1+1)/(5+7) = 2/12 = 1/6$
- $P(\text{Sapporo} | \text{no}) = (2+1)/(5+7) = 3/12 = 1/4$
- $P(\text{Japan} | \text{no}) = (1+1)/(5+7) = 2/12 = 1/6$
- $P(\text{Osaka} | \text{no}) = (1+1)/(5+7) = 2/12 = 1/6$
- $P(\text{Taipei} | \text{no}) = (0+1)/(5+7) = 1/12$
- $P(\text{Macao} | \text{no}) = (0+1)/(5+7) = 1/12$
- $P(\text{Shanghai} | \text{no}) = (0+1)/(5+7) = 1/12$

Part (c): Classify Test Document d5 = {Taiwan, Taiwan, Sapporo}

$\text{Score}(\text{yes}) = \log P(\text{yes}) + \log P(\text{Taiwan}|\text{yes}) + \log P(\text{Taiwan}|\text{yes}) + \log P(\text{Sapporo}|\text{yes})$

- $= \log(0.5) + \log(1/4) + \log(1/4) + \log(1/12)$
- $= -0.693 + (-1.386) + (-1.386) + (-2.485) = -5.950$

$\text{Score}(\text{no}) = \log P(\text{no}) + \log P(\text{Taiwan}|\text{no}) + \log P(\text{Taiwan}|\text{no}) + \log P(\text{Sapporo}|\text{no})$

- $= \log(0.5) + \log(1/6) + \log(1/6) + \log(1/4)$
- $= -0.693 + (-1.792) + (-1.792) + (-1.386) = -5.663$

Result: $\text{Score}(\text{no}) = -5.663 > \text{Score}(\text{yes}) = -5.950$, so d5 is classified as NOT China (no).